

Report: DINOv2 Self-Supervised Learning for MedLeaf Classification

1. Executive Summary of the Code

The notebook implements a computer vision pipeline using **DINOv2** (a self-supervised Vision Transformer) to classify medicinal plant leaf diseases. The workflow utilizes the **MedLeaf dataset**, containing classes such as 'Sojina', 'Neem', 'Camphor', and 'HariTaki' with various disease states (e.g., Healthy, Yellow Leaf, Bacterial Spot).

Key Implementation Steps:

- **Data Preparation:** The code downloads the dataset, resizes images (typically to 518x518 or 224x224 for ViTs), and splits them into training (80%), validation (10%), and testing (10%) sets.
 - **Model Architecture:** It leverages the `timm` (PyTorch Image Models) library to load a pre-trained **DINOv2 ViT (Vision Transformer)** backbone.
 - **Feature Extraction Strategy:** Crucially, the code **freezes the backbone parameters**. This means the massive DINOv2 network is used strictly as a feature extractor. Only a lightweight classification head (Linear layers) is trained on top of these features.
 - **Training:** The model is optimized using standard CrossEntropyLoss and an Adam or SGD optimizer, tracking accuracy metrics across the plant classes.
-

2. How DINOv2 Works (Self-Supervised Learning)

To understand this code, one must understand that DINOv2 is not trained like a standard supervised model (which relies on human labels). It uses **Self-Supervised Learning (SSL)**.

The Mechanism:

1. **No Labels Required:** DINOv2 learns from massive amounts of unlabelled data. It figures out what an image "is" by understanding the relationships between different parts of the image itself.
2. **Student-Teacher Distillation:** The architecture uses two networks with the exact same structure: a **Student** and a **Teacher**.
 - **Global vs. Local Crops:** The model takes an image and creates "global views" (large crops) and "local views" (small patches).
 - **The Task:** The Student sees only the local crops, while the Teacher sees the global views. The Student must try to output the same feature representation as the Teacher. Essentially, the Student learns to guess the "whole picture" from just seeing a small piece of it.
3. **Self-Distillation:** The Teacher network is not updated via standard backpropagation. Instead, it is an **Exponential Moving Average (EMA)** of the Student network. This

creates a stable target for the Student to learn from, preventing the model from "collapsing" (cheating by outputting constant values).

3. Why DINOv2 Yields Better Accuracy

The code likely achieves high accuracy very quickly (even with a simple linear classifier) for several specific reasons inherent to the DINOv2 design:

A. Robust Feature Representations Traditional Convolutional Neural Networks (CNNs) trained with supervision tend to focus on textures or specific discriminative parts of an image (e.g., the ear of a cat). DINOv2, because it must reconstruct global context from local patches, learns a semantic understanding of the **entire object geometry**. It understands the "shape" and "structure" of the leaf, not just the color of the disease spots.

B. The "Freezing" Advantage In the provided code, the backbone is frozen. Because DINOv2 has already "seen" 142 million images during its pre-training, its internal understanding of edges, shapes, and biological structures is already perfect.

- **Data Efficiency:** You do not need thousands of images of "Neem leaves" to train the model. The model already knows what a leaf looks like; you only need to teach the final layer the specific difference between "Healthy" and "Bacterial Spot."
- **Prevention of Overfitting:** By not training the massive backbone, you prevent the model from memorizing the small training set, leading to better generalization on the test set.

C. High Resolution & Patching DINOv2 uses a Vision Transformer (ViT) architecture which breaks images into patches (like a puzzle). This allows it to capture long-range dependencies across the image (e.g., relating a yellow spot on the top left to a withered stem on the bottom right) much better than standard CNNs, which are limited by small receptive fields.